TDDE31/732A54 - Big Data Analytics

Lab compendium
updated 2019-04-09

For relational databases lab, please refer to http://www.ida.liu.se/~732A54/lab/rdb/index.en.shtml.
Notice*: Please make sure you have read the whole lab compendium before you start to work on the server from NSC.*

Description and Aim

In the labs you will work on the Hadoop cluster (Heffa) set up at the National Supercomputer Centre (NSC). You will work with the historical meteorological data from the Swedish Meteorological and Hydrological Institute (SMHI). Specifically, you will work with air temperature readings and precipitation readings from 812 meteorological stations in Sweden[1]. In these exercises, you will work with both Spark and Spark SQL.

After completing the first two labs you will have basic knowledge of the programing environment, techniques and APIs for running both Spark and Spark SQL. You will work on exercises with Spark and Spark SQL and thus will be able to compare the differences between the two approaches. In the third lab, you are supposed to achieve a machine learning method with Spark.

Data

The data includes air temperature and precipitation readings from 812 stations in Sweden. The stations include both currently active stations as well readings from historical stations that have been closed down. The latest readings available for active stations are from October 10, 2016. The air temperature/precipitation readings are hourly readings, however some stations provide only one reading every three hours. The provided files are prepared csv files with removed headers. Values are separated with ;. Some files are too big to be read using some text editors. Therefore, please use either python to read the files or bash commands such as `tail` and `more` to get an overview of a file's content. Provided files:

- *temperature-readings.csv* - ca 2 GB
- *precipitation-readings.csv* - ca 660 MB
- *stations.csv*
- *stations-Ostergotland.csv*
- *temperatures-big.csv* - ca 17 GB

Note: For the relatively big files, they are already available on `hdfs` under: `/user/common/732A54/` and/or locally `heffa` under `/nfshome/hadoop_examples/shared_data/`. For the other files, you can download from https://www.ida.liu.se/~732A54/lab/station_data.zip[2]. If you already log in `heffa` before you download zip file, you just need to upload station files to `hdfs`. Otherwise, you need to upload station files to `heffa` by using `scp` command.

Note: The heffa server has limited space for `/nfshome,` so please don't keep copies of big files in your home folder.

---

[1] If interested in other readings please check: http://opendata-catalog.smhi.se/explore/
[2] To unzip the file, use: `unzip -a station_data.zip`

Headers for *temperature-readings.csv*

Station number	Date	Time	Air temperature (in °C)	Quality[3]

Headers for *precipitation-readings.csv*

Station number	Date	Time	Precipitation (in mm)	Quality[3]

Headers for *stations.csv*

Station number	Station name	Measurement height	Latitude	Longitude	Readings from (date and time)	Readings to (date and time)	Elevation

Headers for *stations-Ostergotland.csv*

These are the same as in stations.csv. The file contains only stations in Östergotland.

Headers for *temperatures-big.csv*

These are the same as in temperature-readings.csv. The file is essentially a concatenation of 8 copies of *temperature-readings.csv* file.

If you notice any mistakes in the dataset/lab compendium or have any comments please contact the lab assistants.

---

[3] G - controlled and confirmed values, Y - suspected or aggregated values

Working on your labs

Cluster setup and logging in

NSC's experimental Heffa lab cluster was built from old nodes from the NSCs 'matter' supercomputer, which was decommissioned. Some details about the nodes are provided below.

System server:	Compute / Login / Analysis nodes:
Hardware: ProLiant DL180 G6 CPU: 2 x 4-core Intel(R) Xeon(R) CPU E5520 @ 2.27GHz Hadoop software: - hadoop hdfs namenode - yarn resource manager - yarn proxyserver - mapreduce historyserver - spark history server	Number of nodes: 11 (of which 2 are login nodes) Hardware: HP SL170z G6 CPUs: 2 x 4-core Intel Xeon E5520 @ 2.2GHz Interconnect: gigabit ethernet Hadoop distributed storage: 9 x 500 GB. Memory: 9 x 4 GB Hadoop software: - hadoop hdfs datanode - hadoop client software (map reduce, etc.) - hadoop yarn nodemanager - spark client software

In the labs you will work with Spark and Spark SQL v. 1.6.0. We will make use of Spark Python API (PySpark) which provides a python programming environment for Spark and Spark SQL. Make use of PySpark's programming guide and API's documentation to get an overview of available functions.[4]

The server is available at `heffa.nsc.liu.se` (log in using your NSC accounts).
The first time you log in after receiving your account details, you must log in using `ssh`. To do this, use the following command in the terminal **from a computer in the lab rooms**:

- `ssh -X` username@heffa.nsc.liu.se where username is your NSC username (not the LiU one), `-X` indicates forwarding function of `ssh` which is used for running graphics applications remotely. When you logout the server, if the logout is hung on, please use ctrl-c to terminate the connection.
- You can use Geany for coding remotely by running `geany &` in the terminal after you connect to heffa or program locally, then use `scp` to copy you files to heffa.
- `[username@heffa1 ~]$ geany &`
- `scp LOCAL_FILE username@heffa.nsc.liu.se:Desktop`

It is always a good practice to verify that one has kerberos tickets before starting to work with Hadoop, and if not, obtain them. You list kerberos tickets by running `klist` in the terminal, and get new ones with `kinit`. An example of a ticket is given below:

```
Default principal: huali50@HEFFA.NSC.LIU.SE

Valid starting     Expires          Service principal
03/20/2018 08:51:26  03/27/2018 09:51:26  krbtgt/HEFFA.NSC.LIU.SE@HEFFA.NSC.LIU.SE
    renew until 04/03/2018 09:51:26
```
Check that you have acquired kerberos **tickets every time before starting your work** with the Hadoop server.

---

[4] http://spark.apache.org/docs/1.6.0/programming-guide.html

Running your scripts

To submit a job to the cluster:

```
spark-submit --deploy-mode cluster --master yarn --num-executors 9 --driver-memory
2g --executor-memory 2g --executor-cores 4 job.py
```
where `job.py` is your python script in your current folder. In this command, we use Yarn for resource management and use the cluster deploy mode. We have 9 worker nodes with 4 cores each with allocated 2GB of memory each.

To make the calling of your python scripts easier, you can download a bash script which includes all the settings (https://www.ida.liu.se/~732A54/lab/scripts/runYarn.sh). In this case, to run your `job.py` you will need to run:

```
./runYarn.sh job.py
```

You can change the settings by editing the `runYarn.sh` file. You might need to add the execute permissions to the script before you run it. To do this run:

```
chmod u+x runYarn.sh
```

During the execution of the job Spark starts SparkUI which is a web user interface for monitoring the job execution (more information available at: http://spark.apache.org/docs/latest/monitoring.html). However, the monitoring will only be available during the execution. In order to be able to access the logs after the execution you will need to set the `spark.eventLog.enabled` flag when running your job:

```
spark-submit --conf spark.eventLog.enabled=true --deploy-mode cluster --master yarn
--num-executors 9 --driver-memory 2g --executor-memory 2g --executor-cores 4 job.py
```

The script which includes the configuration for running the history server is provided here (https://www.ida.liu.se/~732A54/lab/scripts/runYarn-withHistory.sh). To run the job:

```
 ./runYarn-withHistory.sh job.py
```

To access the logs, you can visit http://heffa-head.local:18088 with Firefox. Similar as with runYarn.sh you might need to add the execute permissions.

You can open Firefox by just run it in the terminal.

- `[username@heffa1 ~]$ firefox &`

Scheduling

Given the number of course participants and limited resources it may happen that you experience delays in executing your programs using Yarn. More specifically, you will notice that in some cases your application will be in the ACCEPTED state for few minutes until it reaches the RUNNING state. The reason for this is that there are already running tasks on the cluster which were submitted before. To check the up-to-date information about running/scheduled tasks, you can visit: http://heffa-head.local:8088/cluster through Firefox.

The exercises should not require a lot of time to run, and long running times might imply that there is something wrong with your code. So if you experience long run-times and you do not see other more running jobs please send the application id to the lab assistant for killing the job.

SparkContext

When working with pyspark you will first need to acquire a SparkContext. SparkContext is the entry point to all functionality in Spark. Do this by including the following:

```
from pyspark import SparkContext sc = SparkContext()
```

`SparkContext` accepts a number of parameters, such as the application name, number of executors, etc. For more information, check the documentation. When working with Spark SQL (for BDA2), in addition to SparkContext you will also need to acquire the SQLContext by:

```
from pyspark.sql import SQLContext sqlContext = SQLContext(sc)
```
Where `sc` is your `SparkContext`.

HDFS

In the exercises you will be required to copy files from/to hdfs. In these cases, you will need to make use of `hdfs` commands. Check available commands by running `hdfs dfs` in the terminal. Some useful commands:

`hdfs dfs -ls`   - check the content of the folder

`hdfs dfs -mkdir data`   - make a folder called data

`hdfs dfs -rm file.txt`   - remove the file file.txt

`hdfs dfs -rm -r folder`   - remove the folder and its content

`hdfs dfs -cp /user/common/732A54/temperatures-big.csv /user/username/data` -copy the big file to your data folder on hdfs

`hdfs dfs -copyFromLocal file.txt data/` - copies local file file.txt to folder data on hdfs

`hdfs dfs -copyToLocal results/ .`   - copy the results/ folder to the current folder

When referencing files on hdfs (e.g. with `sc.textFile(path)`) you will need to provide the full path on hdfs. For example, if you created a file `file.txt` under folder data in your home directory on `hdfs`, the full path will be:

```
/user/{username}/data/file.txt
```
where `{username}` is your username.

Reports

For each lab hand in a lab-report following the submission rule. For each exercise, provide your program, results from the program execution (a snippet of the results is enough if the results contain many rows) and written answers to questions in exercises. In cases where a plot of your results is asked, you can include the figure directly in the report. You can use a tool of your preference to produce the plots (e.g. R, Excel, matplotlib in Python, etc.). Comment each step in your program to provide a clear picture of your reasoning when solving the problem.